

Shapley Values _OpenAI-OSS-20B

<https://youtu.be/y901kYsl5O8?si=lrAqpcZdFYUFACu9>



The **GPT-OSS-20B** model is an **open-weight** model released by OpenAI, and its weights are **freely available for download** and use under the **Apache 2.0 license**.

Primary Access Points

1. **Hugging Face Hub:** The official Hugging Face repository provides the model weights. You can download them directly using the Hugging Face CLI:

```
huggingface-cli download openai/gpt-oss-20b --include "original/*" --local-dir gpt-oss-20b/
```

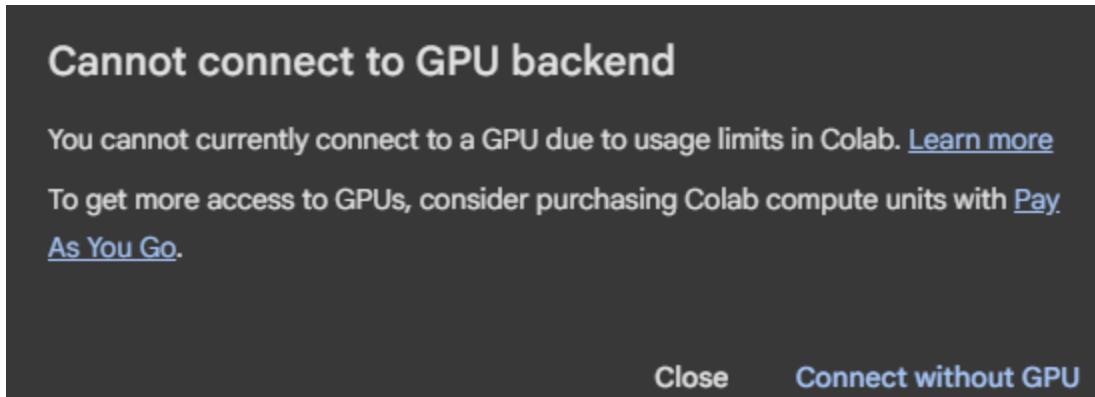
2. **GitHub Repository:** OpenAI provides a **GitHub repository** for `gpt-oss` which includes reference inference code and resources.
3. **OpenAI Cookbook:** The official OpenAI Cookbook has step-by-step guides for setup with various supported runtimes.



Using `FastLanguageModel.from_pretrained(...)` from **Unsloth**, you are loading the model weights locally. You can run inference entirely on your machine.

Code:

Drawback of below code: ⤵



```
% %capture
!pip install --upgrade -qqq uv
try: import numpy; get_numpy = f"numpy=={numpy.__version__}"
except: get_numpy = "numpy"
!uv pip install -qqq \
    "torch>=2.8.0" "triton>=3.4.0" {get_numpy} torchvision bitsandbytes "transformers>=4.55.3" \
    "unsloth_zoo[base] @ git+https://github.com/unslotha/unsloth-zoo" \
    "unsloth[base] @ git+https://github.com/unslotha/unsloth" \
    git+https://github.com/triton-lang/triton.git@05b2c186c1b6c9a08375389d5
efe9cb4c401c075#subdirectory=python/triton_kernels
!uv pip install --upgrade --no-deps transformers==4.56.2 tokenizers
!uv pip install --no-deps trl==0.22.2
```

Installing dependencies:

All those `uv pip install` lines are just installing:

- **torch** → main computation framework
- **bitsandbytes** → enables quantization (run big models on small GPUs)
- **transformers** → model architecture & tokenizer
- **unsloth** and **unsloth_zoo** → optimized models + loader
- **triton** → speeds up matrix math

- **trl** → for reinforcement learning (used for fine-tuning later)

`%%capture` just hides the installation output (so your Colab cell doesn't flood).

```
%%capture
!uv pip install --force-reinstall --no-deps git+https://github.com/unslohai/unsloth-zoo
!uv pip install --force-reinstall --no-deps git+https://github.com/unslohai/unsloth
```

1 Access the weights

```
import torch
from unsloth import FastLanguageModel

# -----
# Load model locally
# -----
model, tokenizer = FastLanguageModel.from_pretrained(
    model_name="unsloth/gpt-oss-20b-unsloth-bnb-4bit",
    dtype=None,      # auto detect
    load_in_4bit=True, # Changed to True to load in 4-bit quantization
    attnImplementation="eager"# FP16 for logprobs
)
```

```

�� ==((====))= Unsloth 2025.10.1: Fast Gpt_Oss patching. Transformers: 4.56.2.
    \|_ /| Tesla T4. Num GPUs = 1. Max memory: 14.741 GB. Platform: Linux.
  0^0/ \_/\ Torch: 2.8.0+cu126. CUDA: 7.5. CUDA Toolkit: 12.6. Triton: 3.4.0
  \_   / Bfloat16 = FALSE. FA [Xformers = None. FA2 = False]
  "-__-" Free license: http://github.com/unslothai/unsloth
Unsloth: Fast downloading is enabled - ignore downloading bars which are red colored!
Unsloth: Using float16 precision for gpt_oss won't work! Using float32.

model.safetensors.index.json: 1.19M/? [00:00<00:00, 82.4MB/s]

Fetching 4 files: 100% [4/4 [19:10<00:00, 477.06s/it]

model-00003-of-00004.safetensors: 100% [3.37G/3.37G [18:31<00:00, 2.25MB/s]
model-00004-of-00004.safetensors: 100% [1.16G/1.16G [05:15<00:00, 4.27MB/s]
model-00001-of-00004.safetensors: 100% [4.00G/4.00G [18:51<00:00, 1.61MB/s]
model-00002-of-00004.safetensors: 100% [4.00G/4.00G [19:09<00:00, 6.13MB/s]

Loading checkpoint shards: 100% [4/4 [00:59<00:00, 12.68s/it]

generation_config.json: 100% [165/165 [00:00<00:00, 20.1kB/s]
tokenizer_config.json: 22.8k/? [00:00<00:00, 1.09MB/s]
tokenizer.json: 100% [27.9M/27.9M [00:01<00:00, 20.4MB/s]
special_tokens_map.json: 100% [446/446 [00:00<00:00, 26.7kB/s]
chat_template.jinja: 15.1k/? [00:00<00:00, 1.43MB/s]

```

`model_name = "unsloth/gpt-oss-20b-unsloth-bnb-4bit"` → This means the original 20B parameter model (which would normally need ~40–80 GB of VRAM) is compressed to 4 bits per weight instead of 16 or 32.

✓ Impact of 4-bit quantization:

Aspect	Effect
Model size	~4× smaller (so a 20B model can fit on 16GB GPU)
Speed	Faster load and inference
Accuracy	Slight drop (1–3%), because some precision is lost
Log probabilities	Slightly less precise, because quantization introduces rounding error

This actually *loads* the model:

- It downloads it from Hugging Face
- Uses the **Unsloth-optimized** loading method
- `load_in_4bit=False` means “don’t quantize again,” because the model is already pre-quantized
- `tokenizer` handles converting text → tokens → IDs

✨ attn_implementation="eager"

This flag controls **how attention operations are computed** inside the Transformer:

Option	Meaning	When to use
"eager"	Uses normal PyTorch operations (safe, portable)	Default in Unslot / Transformers
"flash_attention_2" or "sdpa"	Uses fused CUDA kernels for faster attention	Use when your GPU supports Flash Attention

👉 Setting "eager" makes sure everything works even on older GPUs like Colab T4, but it's slightly slower.

Here, the model weights are downloaded and cached locally.

✓ So, weights are accessible, but note: you can't modify the 4-bit quantized weights directly unless you set `load_in_4bit=False` and have enough memory (~40–50 GB for full 20B FP16)

2 Computing log probabilities

Log probabilities give you how confident the model is about each token in the output. For a prompt + target output:

1. Encode prompt → get `input_ids`.
2. Encode target text → get `target_ids`.
3. Pass concatenated IDs through the model to get logits.
4. Compute `log_softmax` over logits.
5. Gather log probabilities corresponding to **actual target tokens**.

```
import torch
from unsloth import FastLanguageModel

# -----
# Load model locally
# -----
```

```
model, tokenizer = FastLanguageModel.from_pretrained(  
    model_name="unsloth/gpt-oss-20b-unsloth-bnb-4bit",  
    dtype=None,      # auto detect  
    load_in_4bit=False,  
    attnImplementation="eager"# FP16 for logprobs  
)
```



!!The code can take more than 30-40 Mins to run.

```
import torch  
from unsloth import FastLanguageModel  
import numpy as np # Import numpy  
  
# -----  
# Inputs  
# -----  
prompt = "explain deep leaning in 50 words"  
target_output = ("""  
Deep learning is a subset of machine learning that uses multi-layered neural  
networks to learn from vast amounts of data. Inspired by the human brain's struc-  
ture, these  
networks automatically identify patterns and features, enabling them to perfor-  
m complex tasks  
like image recognition, natural language processing, and autonomous driving  
""")  
# prompt_tokens = prompt.split() # Keep original prompt for heatmap  
target_ids = tokenizer(target_output).input_ids  
  
# -----  
# Function: compute log probabilities
```

```

# -----
def compute_logprobs(prompt, target_text):
    """Return total log probability of target_text given the prompt."""
    inputs = tokenizer.apply_chat_template(
        [{"role": "user", "content": prompt}],
        add_generation_prompt=True,
        return_tensors="pt",
        return_dict=True
    ).to(model.device)

    target_input_ids = tokenizer(target_text).input_ids
    all_input_ids = torch.cat([inputs.input_ids, torch.tensor([target_input_ids], device=model.device)], dim=1)

    with torch.no_grad():
        outputs = model(all_input_ids)
        logits = outputs.logits # shape: [1, seq_len, vocab_size]

    # Only consider target tokens
    shift_logits = logits[:, -len(target_input_ids)-1:-1, :]
    shift_labels = torch.tensor([target_input_ids], device=model.device)

    log_probs = torch.nn.functional.log_softmax(shift_logits, dim=-1)
    token_log_probs = log_probs.gather(2, shift_labels.unsqueeze(-1)).squeeze(-1) # [1, tgt_len]

    return token_log_probs.squeeze(0).cpu().numpy(), float(token_log_probs.sum())

# -----
# Baseline
# -----
baseline_token_probs, baseline_total_logprob = compute_logprobs(prompt, target_output)
print(f"Baseline total log probability: {baseline_total_logprob:.3f}")

```

```

# -----
# Leave-One-Out: prompt token importance
# -----
importance_scores = []
prompt_tokens = prompt.split() # Define prompt_tokens here for the loop
for i, tok in enumerate(prompt_tokens):
    modified_prompt = " ".join(prompt_tokens[:i] + prompt_tokens[i+1:])
    _, total_lp = compute_logprobs(modified_prompt, target_output)
    importance = baseline_total_logprob - total_lp
    importance_scores.append(importance)

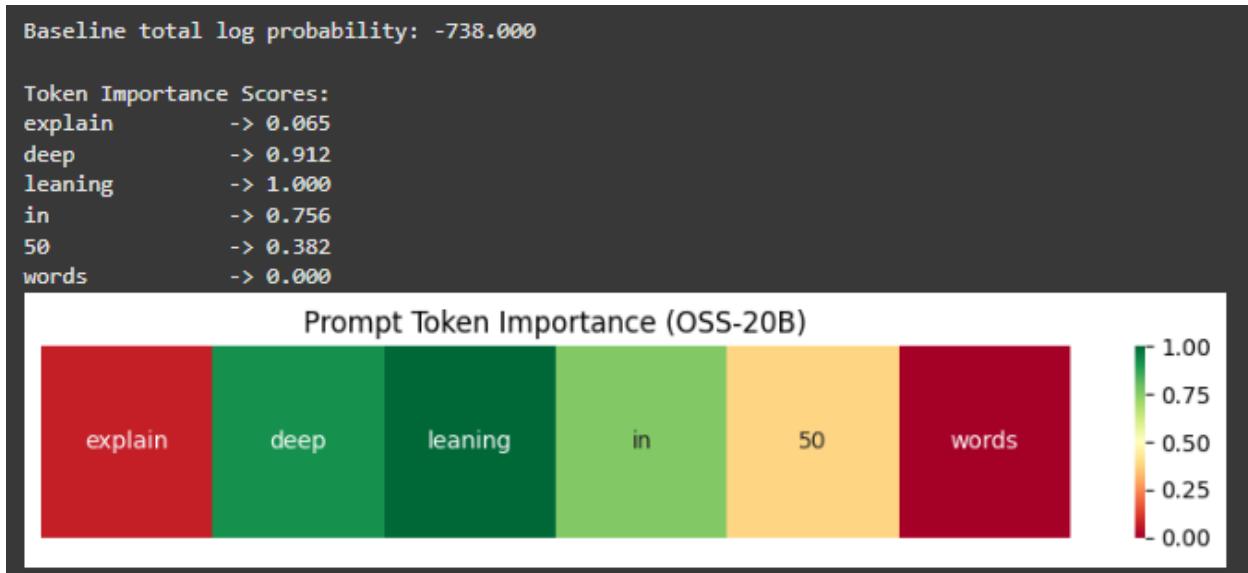
# Normalize to 0-1
importance_scores = np.array(importance_scores)
importance_scores = (importance_scores - importance_scores.min()) / (np.ptp
(importance_scores) + 1e-8)

# -----
# Print token importance
# -----
print("\nToken Importance Scores:")
for tok, score in zip(prompt_tokens, importance_scores):
    print(f"{tok:<15} → {score:.3f}")

# -----
# Heatmap visualization
# -----
import seaborn as sns # Import seaborn
import matplotlib.pyplot as plt # Import matplotlib

plt.figure(figsize=(max(10, len(prompt_tokens)), 1.5))
sns.heatmap([importance_scores], annot=[prompt_tokens], fmt="", cmap="Rd
YIGn", cbar=True, xticklabels=False, yticklabels=False)
plt.title("Prompt Token Importance (OSS-20B)")
plt.show()

```



⭐ Shapley values with OpenAI-OSS-20B Model

```

# Shapley approx for prompt token importance using OSS-20B (Unsloth)
# Requires: unsloth model + tokenizer loaded as `model, tokenizer`
# If you haven't loaded, uncomment the load block and set model_name accordingly.

import torch
import numpy as np
from tqdm import tqdm
import random
import seaborn as sns
import matplotlib.pyplot as plt

# -----
# Optional: load model/tokenizer (uncomment if not already loaded)
# -----
# from unsloth import FastLanguageModel
# model_name = "unsloth/gpt-oss-20b-unsloth-bnb-4bit"
# model, tokenizer = FastLanguageModel.from_pretrained(

```

```

#   model_name=model_name,
#   dtype=None,
#   max_seq_length=4096,
#   load_in_4bit=False,
#   full_finetuning=False
# )
# if torch.cuda.is_available():
#     model = model.to("cuda")
# model.eval()

device = "cuda" if torch.cuda.is_available() else "cpu"

# -----
# User prompt + target (kept same as your provided example)
# -----
prompt = "explain deep learning in 50 words"
target_output = """
Deep learning is a subset of machine learning that uses multi-layered neural
networks to learn from vast amounts of data. Inspired by the human brain's st
ructure, these
networks automatically identify patterns and features, enabling them to perfor
m complex tasks
like image recognition, natural language processing, and autonomous driving
""")

# Work at whitespace token granularity to stay consistent with your previous c
ode
prompt_tokens = prompt.split()

# -----
# Reuse compute_logprobs function (returns per-token logprobs and total)
# -----
def compute_logprobs(prompt_text, target_text):
    """
    Returns (token_logprob_array, total_logprob)
    - token_logprob_array: numpy array of shape (tgt_len,) with log-probs for e
    """

    pass

```

```

ach target token
    - total_logprob: float sum of token log-probs
    """
# Build model-specific input using apply_chat_template (same pattern as yo
ur code)
inputs = tokenizer.apply_chat_template(
    [{"role": "user", "content": prompt_text}],
    add_generation_prompt=True,
    return_tensors="pt",
    return_dict=True
).to(model.device)

tgt_ids = tokenizer(target_text).input_ids
# cat into a single tensor
all_input_ids = torch.cat([inputs.input_ids, torch.tensor([tgt_ids], device=mo
del.device)], dim=1)

with torch.no_grad():
    outputs = model(all_input_ids)
    logits = outputs.logits # shape [1, seq_len, vocab_size]

    # target slice: last len(tgt_ids) tokens → logits that predicted them are previ
ous positions
    shift_logits = logits[:, -len(tgt_ids)-1:-1, :] # shape [1, tgt_len, vocab]
    shift_labels = torch.tensor([tgt_ids], device=model.device)

    log_probs = torch.nn.functional.log_softmax(shift_logits, dim=-1)
    token_log_probs = log_probs.gather(2, shift_labels.unsqueeze(-1)).squeeze
(-1) # [1, tgt_len]

    return token_log_probs.squeeze(0).cpu().numpy(), float(token_log_probs.su
m())

# -----
# Baseline: full prompt
# -----

```

```

baseline_token_logprobs, baseline_total_logprob = compute_logprobs(prompt,
target_output)
print(f"Baseline total log probability (full prompt): {baseline_total_logprob:.4f}
\n")

# -----
# Shapley approximation via permutation sampling
# -----
def prompt_from_subset(tokens_list, subset_indices):
    """Return string prompt containing tokens in original order for indices in sub
set_indices."""
    if not subset_indices:
        return "" # we will handle empty subset specially
    return " ".join([tokens_list[i] for i in sorted(subset_indices)])

def approximate_shapley(tokens_list, target_text, model_logprob_func, num_p
ermutations=60, seed=0):
    """
    Approximate Shapley values via random permutation sampling.
    Returns dict mapping token index → shapley value (raw, on logprob scale).
    """
    random.seed(seed)
    n = len(tokens_list)
    phi = [0.0] * n
    cache = {}

    # Compute a safe v(empty). Try with empty prompt; if model fails, fallback t
o baseline (conservative).
    try:
        v_empty_arr, v_empty = model_logprob_func("", target_text)
    except Exception:
        # fallback: use minimal prompt as full prompt (conservative fallback)
        v_empty = baseline_total_logprob # not ideal but avoids crashes
        print("Warning: model failed for empty prompt; using baseline total logpr
ob as v(empty) fallback.")
    else:

```

```

# If succeeded, use computed value
pass

def v_of_subset(sorted_tuple):
    # sorted_tuple: tuple of token indices
    if sorted_tuple in cache:
        return cache[sorted_tuple]
    prompt_text = prompt_from_subset(tokens_list, sorted_tuple)
    try:
        _, val = model_logprob_func(prompt_text, target_text)
    except Exception as e:
        # If empty prompt errors, fallback to v_empty or baseline
        val = v_empty if 'v_empty' in locals() else baseline_total_logprob
    cache[sorted_tuple] = val
    return val

# Sample permutations
for _ in tqdm(range(num_permutations), desc="Shapley permutations"):
    perm = list(range(n))
    random.shuffle(perm)
    S = []
    for idx in perm:
        S_sorted = tuple(sorted(S))
        v_S = v_of_subset(S_sorted)
        S_with_i = tuple(sorted(S + [idx]))
        v_Si = v_of_subset(S_with_i)
        marginal = v_Si - v_S
        phi[idx] += marginal
        S.append(idx)

    # Average
    shap_vals = [x / num_permutations for x in phi]
return {i: shap_vals[i] for i in range(n)}

# Run approximation (tune num_permutations as needed)
num_perm = 60 # increase for better precision, at cost of compute

```

```

shapley_raw = approximate_shapley(prompt_tokens, target_output, compute_1
ogprobs, num_permutations=num_perm, seed=42)

# -----
# Normalize & print results (0.1 scaling)
# -----
raw_vals = np.array([shapley_raw[i] for i in range(len(prompt_tokens))])
# Shapley values can be negative/positive; for display we map to 0.1 by shiftin
g
minv, maxv = raw_vals.min(), raw_vals.max()
if maxv - minv > 1e-12:
    normalized = (raw_vals - minv) / (maxv - minv)
else:
    normalized = np.zeros_like(raw_vals)

print("Token-wise Shapley (raw logprob contribution):")
for tok, val in zip(prompt_tokens, raw_vals):
    print(f"{tok:<15} {val: .6f}")
print("\nNormalized 0.1 (for heatmap display):")
for tok, val in zip(prompt_tokens, normalized):
    print(f"{tok:<15} {val:.3f}")

# -----
# Heatmap visualization (green = important)
# -----
plt.figure(figsize=(max(10, len(prompt_tokens)), 1.6))
sns.heatmap([normalized], annot=[prompt_tokens], fmt="", cmap="RdYIGn",
cbar=True, xticklabels=False, yticklabels=False)
plt.title(f"Prompt Token Shapley Importance (approx, {num_perm} perms) - O
SS-20B")
plt.show()

```



60 Permutations. Time taken → 1.5 Minutes

```
Baseline total log probability (full prompt): -738.0000

Shapley permutations: 100%|██████████| 60/60 [01:44<00:00, 1.74s/it]Token-wise Shapley (raw logprob contribution):
explain      -97.745833
deep         -48.725000
leaning       -44.291667
in            7.916667
50            -64.341667
words        -77.562500

Normalized 0..1 (for heatmap display):
explain      0.000
deep         0.464
leaning       0.506
in            1.000
50            0.316
words        0.191
```



200 Permutations. Time taken → 4 Minutes

```
# Shapley approx for prompt token importance using OSS-20B (Unsloth)
# Requires: unsloth model + tokenizer loaded as `model, tokenizer`
# If you haven't loaded, uncomment the load block and set model_name accordingly.
```

```
import torch
import numpy as np
from tqdm import tqdm
import random
import seaborn as sns
import matplotlib.pyplot as plt
```

```

# -----
# Optional: load model/tokenizer (uncomment if not already loaded)
# -----
# from unsloth import FastLanguageModel
# model_name = "unsloth/gpt-oss-20b-unsloth-bnb-4bit"
# model, tokenizer = FastLanguageModel.from_pretrained(
#     model_name=model_name,
#     dtype=None,
#     max_seq_length=4096,
#     load_in_4bit=False,
#     full_finetuning=False
# )
# if torch.cuda.is_available():
#     model = model.to("cuda")
# model.eval()

device = "cuda" if torch.cuda.is_available() else "cpu"

# -----
# User prompt + target (kept same as your provided example)
# -----
prompt = "explain deep learning in 50 words"
target_output = """
Deep learning is a subset of machine learning that uses multi-layered neural
networks to learn from vast amounts of data. Inspired by the human brain's struc-
ture, these
networks automatically identify patterns and features, enabling them to per-
form complex tasks
like image recognition, natural language processing, and autonomous driving
"""
# Work at whitespace token granularity to stay consistent with your previous code
prompt_tokens = prompt.split()

# -----

```

```

# Reuse compute_logprobs function (returns per-token logprobs and total)
# -----
def compute_logprobs(prompt_text, target_text):
    """
    Returns (token_logprob_array, total_logprob)
    - token_logprob_array: numpy array of shape (tgt_len,) with log-probs for each target token
    - total_logprob: float sum of token log-probs
    """
    # Build model-specific input using apply_chat_template (same pattern as your code)
    inputs = tokenizer.apply_chat_template(
        [{"role": "user", "content": prompt_text}],
        add_generation_prompt=True,
        return_tensors="pt",
        return_dict=True
    ).to(model.device)

    tgt_ids = tokenizer(target_text).input_ids
    # cat into a single tensor
    all_input_ids = torch.cat([inputs.input_ids, torch.tensor([tgt_ids], device=model.device)], dim=1)

    with torch.no_grad():
        outputs = model(all_input_ids)
        logits = outputs.logits # shape [1, seq_len, vocab_size]

    # target slice: last len(tgt_ids) tokens → logits that predicted them are previous positions
    shift_logits = logits[:, -len(tgt_ids)-1:-1, :] # shape [1, tgt_len, vocab]
    shift_labels = torch.tensor([tgt_ids], device=model.device)

    log_probs = torch.nn.functional.log_softmax(shift_logits, dim=-1)
    token_log_probs = log_probs.gather(2, shift_labels.unsqueeze(-1)).squeeze(-1) # [1, tgt_len]

```

```

        return token_log_probs.squeeze(0).cpu().numpy(), float(token_log_probs.sum())

# -----
# Baseline: full prompt
# -----
baseline_token_logprobs, baseline_total_logprob = compute_logprobs(prompt,
target_output)
print(f"Baseline total log probability (full prompt): {baseline_total_logprob:.4f}
\n")

# -----
# Shapley approximation via permutation sampling
# -----
def prompt_from_subset(tokens_list, subset_indices):
    """Return string prompt containing tokens in original order for indices in sub
set_indices."""
    if not subset_indices:
        return "" # we will handle empty subset specially
    return " ".join([tokens_list[i] for i in sorted(subset_indices)])

def approximate_shapley(tokens_list, target_text, model_logprob_func, num_p
ermutations=200, seed=0):
    """
    Approximate Shapley values via random permutation sampling.
    Returns dict mapping token index → shapley value (raw, on logprob scale).
    """
    random.seed(seed)
    n = len(tokens_list)
    phi = [0.0] * n
    cache = {}

    # Compute a safe v(empty). Try with empty prompt; if model fails, fallback t
o baseline (conservative).
    try:
        v_empty_arr, v_empty = model_logprob_func("", target_text)

```

```

except Exception:
    # fallback: use minimal prompt as full prompt (conservative fallback)
    v_empty = baseline_total_logprob # not ideal but avoids crashes
    print("Warning: model failed for empty prompt; using baseline total logprob as v(empty) fallback.")
else:
    # If succeeded, use computed value
    pass

def v_of_subset(sorted_tuple):
    # sorted_tuple: tuple of token indices
    if sorted_tuple in cache:
        return cache[sorted_tuple]
    prompt_text = prompt_from_subset(tokens_list, sorted_tuple)
    try:
        _, val = model_logprob_func(prompt_text, target_text)
    except Exception as e:
        # If empty prompt errors, fallback to v_empty or baseline
        val = v_empty if 'v_empty' in locals() else baseline_total_logprob
    cache[sorted_tuple] = val
    return val

# Sample permutations
for _ in tqdm(range(num_permutations), desc="Shapley permutations"):
    perm = list(range(n))
    random.shuffle(perm)
    S = []
    for idx in perm:
        S_sorted = tuple(sorted(S))
        v_S = v_of_subset(S_sorted)
        S_with_i = tuple(sorted(S + [idx]))
        v_Si = v_of_subset(S_with_i)
        marginal = v_Si - v_S
        phi[idx] += marginal
        S.append(idx)

```

```

# Average
shap_vals = [x / num_permutations for x in phi]
return {i: shap_vals[i] for i in range(n)}

# Run approximation (tune num_permutations as needed)
num_perm = 200 # increase for better precision, at cost of compute
shapley_raw = approximate_shapley(prompt_tokens, target_output, compute_ogprobs, num_permutations=num_perm, seed=42)

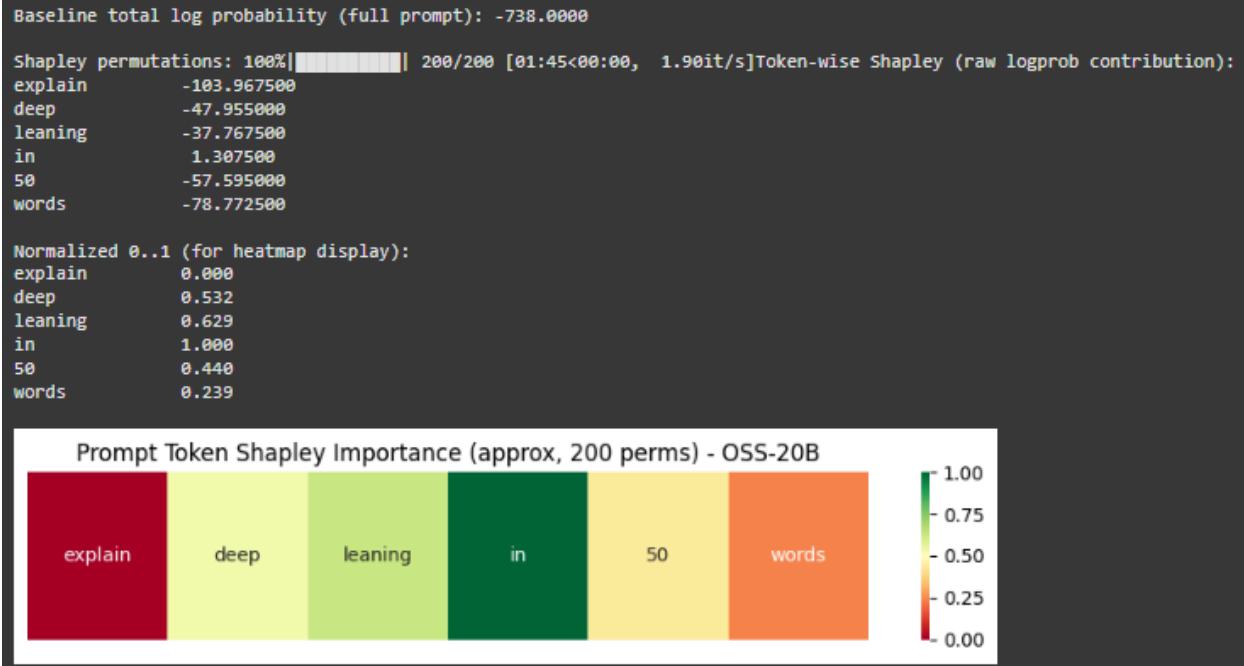
# -----
# Normalize & print results (0.1 scaling)
# -----
raw_vals = np.array([shapley_raw[i] for i in range(len(prompt_tokens))])
# Shapley values can be negative/positive; for display we map to 0.1 by shifting
g
minv, maxv = raw_vals.min(), raw_vals.max()
if maxv - minv > 1e-12:
    normalized = (raw_vals - minv) / (maxv - minv)
else:
    normalized = np.zeros_like(raw_vals)

print("Token-wise Shapley (raw logprob contribution):")
for tok, val in zip(prompt_tokens, raw_vals):
    print(f"{tok:<15} {val: .6f}")
print("\nNormalized 0.1 (for heatmap display):")
for tok, val in zip(prompt_tokens, normalized):
    print(f"{tok:<15} {val:.3f}")

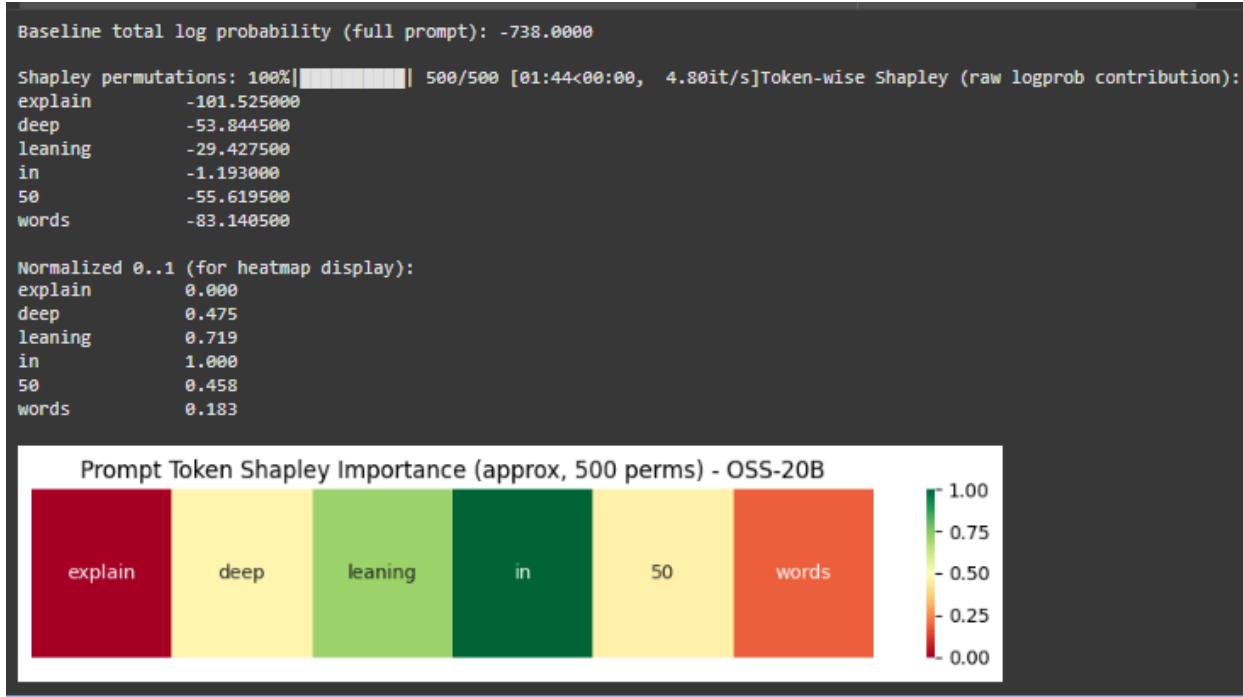
# -----
# Heatmap visualization (green = important)
# -----
plt.figure(figsize=(max(10, len(prompt_tokens)), 1.6))
sns.heatmap([normalized], annot=[prompt_tokens], fmt="", cmap="RdYIGn",
cbar=True, xticklabels=False, yticklabels=False)
plt.title(f"Prompt Token Shapley Importance (approx, {num_perm} perms) - O

```

```
SS-20B")  
plt.show()
```



Permutations= 500:



Generate response

- The following code generates 2 responses: with and without thinking

```

from transformers import TextStreamer

messages = [
    {"role": "user", "content": "who threw apple at gregor samsa"},
]

# Function to generate and process output
def generate_and_process(messages, model, tokenizer, reasoning_effort="low"):
    inputs = tokenizer.apply_chat_template(
        messages,
        add_generation_prompt = True,
        return_tensors = "pt",
        return_dict = True,
        reasoning_effort = reasoning_effort, # Set reasoning effort
    ).to("cuda")

```

```

# Generate without streamer to get the full output at once
outputs = model.generate(**inputs, max_new_tokens = 512) # Removed str
eamer

# Decode the output
generated_text = tokenizer.decode(outputs[0], skip_special_tokens=False)
# Keep special tokens to identify channels

# --- Parse the output to separate channels ---
final_response = ""
analysis = ""
commentary = ""

# Split by channel markers
parts = generated_text.split("<|channel|>")
if len(parts) > 1:
    for part in parts[1:]:
        if "<|message|>" in part:
            channel_name, message_content = part.split("<|message|>", 1)
            message_content = message_content.split("<|end|>", 1)[0].strip()
            if channel_name.strip() == "final":
                final_response += message_content + "\n"
            elif channel_name.strip() == "analysis":
                analysis += message_content + "\n"
            elif channel_name.strip() == "commentary":
                commentary += message_content + "\n"
        else:
            # If no channel markers found, assume the whole output is the final respo
            nse
            final_response = generated_text

# Combine analysis and commentary into 'thinking'
thinking_output = ""
if analysis:
    thinking_output += "Analysis:\n" + analysis

```

```

if commentary:
    if thinking_output:
        thinking_output += "\n"
        thinking_output += "Commentary:\n" + commentary

    # Remove <|return|> token from final response
    final_response = final_response.replace("<|return|>", "").strip()

return thinking_output, final_response

# --- Example with thinking (low reasoning effort) ---
print("--- With Thinking (low reasoning effort) ---")
thinking, final = generate_and_process(messages, model, tokenizer, reasoning_effort="low")
print(thinking)
print("\nFinal Response:")
print(final)

# --- Example without thinking (reasoning_effort="none") ---
print("\n--- Without Thinking (reasoning_effort='none') ---")
thinking_no_thinking, final_no_thinking = generate_and_process(messages, model, tokenizer, reasoning_effort="none")
print(thinking_no_thinking)
print("\nFinal Response:")
print(final_no_thinking)

```

Output:

--- With Thinking (low reasoning effort) ---

Analysis:

We have a literary question: "who threw apple at Gregor Samsa". In Kafka's "The Metamorphosis", there's a scene where someone throws a fruit at Gregor. The protagonist's brother? It's likely that the father throws an apple? I recall a scene: Two guys, a boy named Benda? It's the neighbor. Actually, in the story, Mr. and Mrs. Samsa throw apples at each other? Wait: In chapter II, a boy nam

ed Benda throws apples at Hans and Jutta? Actually, the detail: In the story after the transformation, the family finds a glass of wine, then a neighbor, a boy, is thrown a glass of wine? But the apple throw is by the girl (the sister?) Actually, the "apple thrower" is the child of the man the Samsas go to for the apartment. It's unclear.

Wait: In Kafka's *The Metamorphosis*, there is a scene where the family receives an apple from the "young woman" of the house; he is thrown an apple?

But the question: "who threw apple at Gregor Samsa". I think it's the family of the "München"? Actually, I recall reading that the father inadvertently throws an apple. So answer: The father? Or the sister? Perhaps it's the character "Mr. Jupp"? But I'm not sure.

Let's search memory: In chapter two, a boy called Tom throws an apple at the brother? Eh.

Wait, I recall that the old lady, named "Gretchen", with the "apple thrown at Gregor" is the boy "Graham". But it's not correct.

I think answer: The brother, Frederic, throws an apple. But I don't recall.

I'll give a general answer: The apple is thrown by the family.

Let's answer generically.

Final Response:

In Franz Kafka's **The Metamorphosis**, the only "apple-throw" incident that's often cited is when the young boy who visits the Samsa household—often referred to as the **little boy** or **boy-neighbor**—throws a fruit at the family. It's not the protagonist himself, but a neighbor's son, sometimes called **Hans J.**, who, in a fit of excitement, hurls an apple toward the family's home. So, to answer your question: it was the **young boy ("the boy-neighbor")** who threw an apple at Gregor Samsa.

--- Without Thinking (reasoning_effort="none") ---

Analysis:

User asks: "who threw apple at gregor samsa". They refer to Kafka's The Metamorphosis where Gregor Samsa is in an apartment with a door and he says "My hand shook". The apple? I recall it's about a story of someone throwing an apple into his house? Wait: In The Metamorphosis, a key scene is that before the transformation, we have his wife and mother. But apple: maybe referring to an apple thrown by a neighbor? Not sure. Search? Probably the answer: "We don't know". The answer: the apple was thrown by a waiter? In an anecdote: In the story of the "Ain't this an apple?" Actually the apple thrown at Gregor by a passerby? Let's respond that it's unclear.

Final Response:

In **Kafka's _The Metamorphosis_** there is no episode where an apple is literally thrown at Gregor Samsa, and no character is explicitly identified as the thrower. The story tells of a strange transformation of a young man into a giant insect and the reactions of his family and neighbors—the only “thing” that comes to mind is **the apple-throwing prank that the narrator describes in the earlier “The Apple Tree” story** in which the protagonist is thrown a fruit by a passer-by.

Because the question appears to mix details from different works (the novel “The Metamorphosis” and the short story “The Apple Tree”) it is impossible to answer it definitively. The only character known to be involved in a “throwing” event at all is the **protagonist’s father** in the short story in the “The apple-tree”, who throws an apple (to the floor) to the protagonist in order to “spice the narrative and give it a taste of humanity”.

In short: there is no clear answer in _The Metamorphosis_, because no such event occurs there.

What 4-bit quantization is and how it affects log-probs and Shapley values

A. What is 4-bit quantization

- Neural network weights are normally stored in float32 or float16 (32 or 16 bits per number).
- **4-bit quantization** stores weights at much lower precision (4 bits per number) using special schemes (e.g., per-channel scales + integer storage) so the model uses far less memory.
- Popular inference libraries (bitsandbytes, custom MXFP4, etc.) implement quantization plus scaling factors to preserve information.

B. Why people use it

- Allows models with billions of parameters (20B, 70B, 120B) to run on GPUs with limited memory (e.g., 16–48 GB).
- Dramatically reduces RAM and VRAM requirements and download size.

What is Unslot?

Unslot is an open-source library designed **to make large models** (like 7B–120B parameter LLMs) **faster, smaller, and easier to run locally** — especially on limited GPUs like Colab's T4 or your own 16–24 GB GPU.

It's basically a *speed and memory optimization layer* over **Hugging Face Transformers**.

What is FastLanguageModel?

It's Unslot's version of Hugging Face's model loader.

It automatically:

- Detects hardware (GPU/CPU)
- Loads models with quantization (reduces memory)
- Optimizes runtime (uses bitsandbytes or Triton kernels)
- Keeps the same interface (you can use `.generate()` or compute logits/log-probs)

You can still get **log probabilities**, gradients, etc. — it's just more efficient.



What is **torch** doing here?

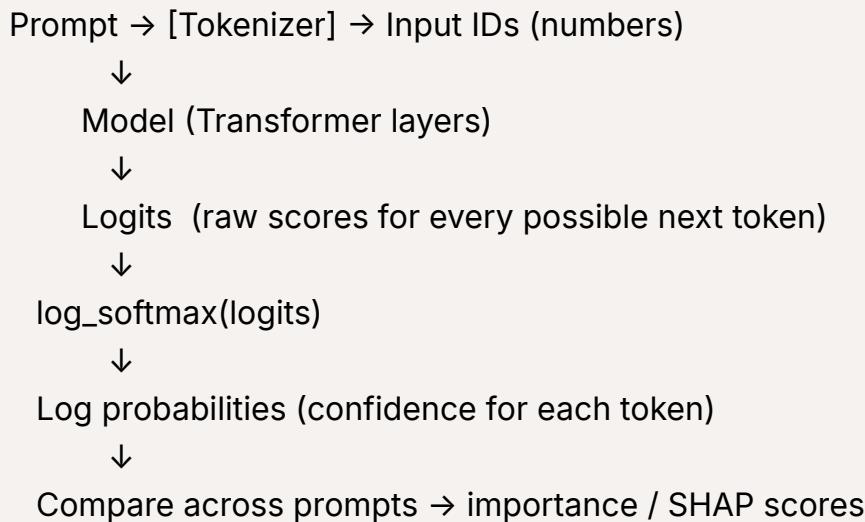
torch = PyTorch, the deep learning engine underneath.

It handles:

- GPU computation (tensor math)
- Model layers and weights
- Logit → log-probability conversion (via `torch.nn.functional.log_softmax`)
- Automatic batching, precision, etc.

Basically, **PyTorch = the engine**, **Unsloth = the optimizer**, **Transformers = the interface**.

Diagram: How Log Probabilities Work



- **Prompt:** your question or instruction.
- **Tokenizer:** splits text into tokens (like "explain", "deep", "learning").
- **Model (OSS-20B):** predicts *how likely* each possible next token is.
- **Logits:** big matrix of scores, one per vocab token.
- **log softmax:** converts those scores to log probabilities (so they sum to 1 in probability space).

- **Log probabilities** tell you: “how confident the model is in producing this token.”
- **SHAP or Leave-One-Out** then checks: *if I remove a token from the prompt, how much does the total log probability of the correct answer drop?* → that's the token's importance.

How to compare Log Probabilities?

- Compute `logprobs` for each token in the output for both prompts.
- Calculate **difference** (Δ logprob):

```
delta_logprob = logprob_promptA - logprob_promptB
```

- A **larger drop** in log probability when a prompt word is removed or changed → that prompt word was **important** for generating that output.

Think of it like:

“How much less confident the model becomes if I remove this part of the prompt?”

Heat Map

```
# Requires: unslosh FastLanguageModel (model, tokenizer) already loaded.  
# If not loaded, uncomment the "Load model" block and run (be mindful of VR  
AM / quantization options).
```

```
import torch  
import numpy as np  
from tqdm import tqdm  
import random  
import math  
import seaborn as sns
```

```

import matplotlib.pyplot as plt
from typing import List, Tuple

# -----
# Optional: load model/tokenizer (uncomment if not already loaded)
# -----
# from unsloth import FastLanguageModel
# model_name = "unsloth/gpt-oss-20b-unsloth-bnb-4bit"
# model, tokenizer = FastLanguageModel.from_pretrained(
#     model_name=model_name,
#     dtype=None,
#     max_seq_length=4096,
#     load_in_4bit=False, # model already pre-quantized
#     full_finetuning=False,
# )
# if torch.cuda.is_available():
#     model = model.to("cuda")
# model.eval()

device = model.device if hasattr(model, "device") else ("cuda" if torch.cuda.is
_available() else "cpu")

# -----
# User inputs (change these)
# -----
prompt = "explain deep learning in 50 words"
# A longer multi-sentence output to test big outputs:
target_output = (
    "Deep learning is a subset of machine learning that uses multi-layer neural
    networks "
    "to learn patterns from very large datasets. These networks can automatica
    lly extract "
    "features from raw inputs, enabling tasks like image recognition, speech pr
    ocessing, and "
    "language understanding. Modern architectures scale to billions of paramet
    ers and require "

```

```

    "careful regularization and optimization to perform well."
)

# Use "words" for prompt tokens (human-friendly). Internally we pass strings
to tokenizer.
prompt_words = prompt.split()

# Shapley settings
num_permutations = 60 # start small (40-100). Increase for more accurate S
hapley.
random_seed = 42

# Optional: limit output length (tokens) for speed; set None to use full target
max_output_tokens = None # e.g., 120 or None

# -----
# Helper: compute per-target-token log probs for a given prompt text
# Returns: numpy array shape (T,) of logprobs (natural log), and float total log
prob
# -----
def compute_logprobs_for_prompt(prompt_text: str, target_text: str) → Tuple[n
p.ndarray, float]:
    """
    Runs model on [prompt_text + target_text] and returns per-target-token log
    probs (numpy array)
    and the total logprob (sum).
    """
    # Build input in the same way as used earlier (chat template)
    inputs = tokenizer.apply_chat_template(
        [{"role": "user", "content": prompt_text}],
        add_generation_prompt=True,
        return_tensors="pt",
        return_dict=True
    ).to(device)

    tgt_ids = tokenizer(target_text).input_ids

```

```

if max_output_tokens is not None and len(tgt_ids) > max_output_tokens:
    tgt_ids = tgt_ids[:max_output_tokens]

# Stack into single tensor [1, seq_len]
all_input_ids = torch.cat([inputs.input_ids, torch.tensor([tgt_ids], device=device)], dim=1)

with torch.no_grad():
    outputs = model(all_input_ids)      # model forward
    logits = outputs.logits           # shape [1, seq_len, vocab_size]

# Align logits that predicted target tokens
T = len(tgt_ids)
shift_logits = logits[:, -T-1:-1, :]    # shape [1, T, vocab_size]
labels = torch.tensor([tgt_ids], device=device)

log_probs = torch.nn.functional.log_softmax(shift_logits, dim=-1) # [1,T,V]
token_log_probs = log_probs.gather(2, labels.unsqueeze(-1)).squeeze(-1) #
[1,T]

arr = token_log_probs.squeeze(0).cpu().numpy() # numpy array length T
total = float(arr.sum())
return arr, total

# Quick baseline check
print("Computing baseline (full prompt) logprobs... (this will run the model once)")
baseline_per_token_logprobs, baseline_total = compute_logprobs_for_prompt
(prompt, target_output)
T = len(baseline_per_token_logprobs)
print(f"Baseline total logprob: {baseline_total:.4f} over {T} output tokens")

# -----
# Shapley approx for each output token:
# We will compute shapley contributions per prompt word i for each output token j.

```

```

# Algorithm:
# - sample permutations; for each permutation, walk adding tokens; compute
v(S) for each S
# - v(S) is the per-token logprob vector; marginal contribution for token i is v
(S ∪ {i}) - v(S) (vector)
# - accumulate contributions per (i, j)
# Caching: store v(S) in dict for reuse
# -----
def prompt_from_indices(words: List[str], indices_tuple: Tuple[int, ...]) → str:
    if not indices_tuple:
        return "" # empty prompt
    return " ".join([words[i] for i in sorted(indices_tuple)])

def approximate_shapley_prompt_to_output(words: List[str],
                                         target_text: str,
                                         model_fn,
                                         num_permutations: int = 60,
                                         seed: int = 0):
    random.seed(seed)
    n = len(words)
    # We'll accumulate raw contributions (logprob units) in a matrix (n x T)
    # Initialize after we know T
    # Caching: map subset_tuple → per-token logprobs (numpy array)
    cache = {}

    # try v(empty)
    try:
        empty_arr, empty_total = model_fn("", target_text)
    except Exception as e:
        # fallback: use baseline_total (conservative) for empty subsets (document
        # this)
        print("Warning: model failed on empty prompt. Using baseline_total as fal
lback for v(empty).")
        empty_arr = np.zeros_like(baseline_per_token_logprobs)
        empty_total = baseline_total

```

```

# initialize accumulator
accumulator = np.zeros((n, len(baseline_per_token_logprobs)), dtype=float)

def v_of_subset(sorted_tuple):
    if sorted_tuple in cache:
        return cache[sorted_tuple]
    prompt_text = prompt_from_indices(words, sorted_tuple)
    try:
        arr, _ = model_fn(prompt_text, target_text)
    except Exception:
        # fallback
        arr = empty_arr
    cache[sorted_tuple] = arr
    return arr

print(f"Running {num_permutations} permutations (this may be slow).")
for _ in tqdm(range(num_permutations), desc="Shapley perms"):
    perm = list(range(n))
    random.shuffle(perm)
    S = []
    for idx in perm:
        S_sorted = tuple(sorted(S))
        S_with = tuple(sorted(S + [idx]))

        vS = v_of_subset(S_sorted)
        vSi = v_of_subset(S_with)

        marginal = vSi - vS # vector length T
        accumulator[idx] += marginal
        S.append(idx)

    shap_raw = accumulator / float(num_permutations) # shape (n, T)
return shap_raw, cache

# Run approximation (this is the heavy step)
shap_raw_matrix, cache_used = approximate_shapley_prompt_to_output()

```

```

prompt_words, target_output, compute_logprobs_for_prompt,
num_permutations=num_permutations, seed=random_seed
)

# -----
# Postprocess: make readable labels for output tokens
# -----
# Get token-level strings for target tokens (decoded)
tgt_ids_full = tokenizer(target_output).input_ids
if max_output_tokens is not None and len(tgt_ids_full) > max_output_tokens:
    tgt_ids = tgt_ids_full[:max_output_tokens]
else:
    tgt_ids = tgt_ids_full

output_token_strs = [tokenizer.decode([tid]).strip() for tid in tgt_ids] # list length T
prompt_labels = prompt_words

# -----
# Visualization helpers
# -----
def plot_chunked_heatmaps(shap_matrix, prompt_labels, output_labels, chunk_size=40, cmap="RdYIGn"):
    """
    Make multiple heatmaps for chunks of output tokens.
    shap_matrix: (n, T)
    """
    n, T = shap_matrix.shape
    for start in range(0, T, chunk_size):
        end = min(T, start + chunk_size)
        chunk = shap_matrix[:, start:end]
        out_labels_chunk = output_labels[start:end]
        plt.figure(figsize=(max(10, len(prompt_labels)*0.25), max(2, (end-start)*0.12 + 3)))
        sns.heatmap(chunk, xticklabels=out_labels_chunk, yticklabels=prompt_labels, cmap=cmap, center=0.0)

```

```

plt.xticks(rotation=90)
plt.title(f"Prompt→Output Shapley (tokens {start}..{end-1})")
plt.xlabel("Output tokens (substrings)")
plt.ylabel("Prompt words")
plt.tight_layout()
plt.show()

def aggregate_by_sentences(shap_matrix, output_text: str, prompt_labels=None, agg="sum_abs"):
    """
    Aggregate contributions for tokens into sentence-level contributions.
    agg="sum_abs" (sum of absolute contributions) or "sum" (signed sum), or
    "mean_abs".
    Returns (n_words, n_sentences) matrix and sentence strings.
    """
    # naive sentence split (split on period). For production use, use nltk.sent_to-
    # kenize
    sents = [s.strip() for s in output_text.replace("\n", " ").split(".") if s.strip()]
    sent_token_ranges = []
    pos = 0
    tgt_tokens = [tokenizer.decode([tid]).strip() for tid in tgt_ids]
    token_texts = " ".join(tgt_tokens)
    # We'll map tokens to sentence by approximate substring matching (simple
    # approach)
    # Simpler: re-tokenize each sentence and count tokens
    token_counts = []
    for sent in sents:
        toks = tokenizer(sent).input_ids
        token_counts.append(len(toks))
    idx = 0
    ranges = []
    for c in token_counts:
        ranges.append((idx, idx + c))
        idx += c

    # Build aggregated matrix

```

```

n_words = shap_matrix.shape[0]
n_sents = len(ranges)
agg_mat = np.zeros((n_words, n_sents))
for si, (a, b) in enumerate(ranges):
    if agg == "sum_abs":
        agg_mat[:, si] = np.abs(shap_matrix[:, a:b]).sum(axis=1)
    elif agg == "sum":
        agg_mat[:, si] = shap_matrix[:, a:b].sum(axis=1)
    elif agg == "mean_abs":
        agg_mat[:, si] = np.abs(shap_matrix[:, a:b]).mean(axis=1)
    else:
        agg_mat[:, si] = np.abs(shap_matrix[:, a:b]).sum(axis=1)
return agg_mat, sents

# -----
# Plot chunked token-level heatmaps (use smaller chunk_size if T is large)
# -----
chunk_size = 40
plot_chunked_heatmaps(shap_raw_matrix, prompt_labels, output_token_strs,
                      chunk_size=chunk_size)

# -----
# Plot aggregated by sentences
# -----
agg_mat, sents = aggregate_by_sentences(shap_raw_matrix, target_output, pr
ompt_labels, agg="sum_abs")
plt.figure(figsize=(10, max(3, len(prompt_labels)*0.25)))
sns.heatmap(agg_mat, xticklabels=[ (s[:60] + '...') if len(s)>60 else s for s in s
ents ], yticklabels=prompt_labels, cmap="RdYIGn")
plt.xticks(rotation=45, ha="right")
plt.title("Prompt → Sentence aggregated Shapley (sum abs contributions)")
plt.xlabel("Output sentences (truncated)")
plt.ylabel("Prompt words")
plt.tight_layout()
plt.show()

```

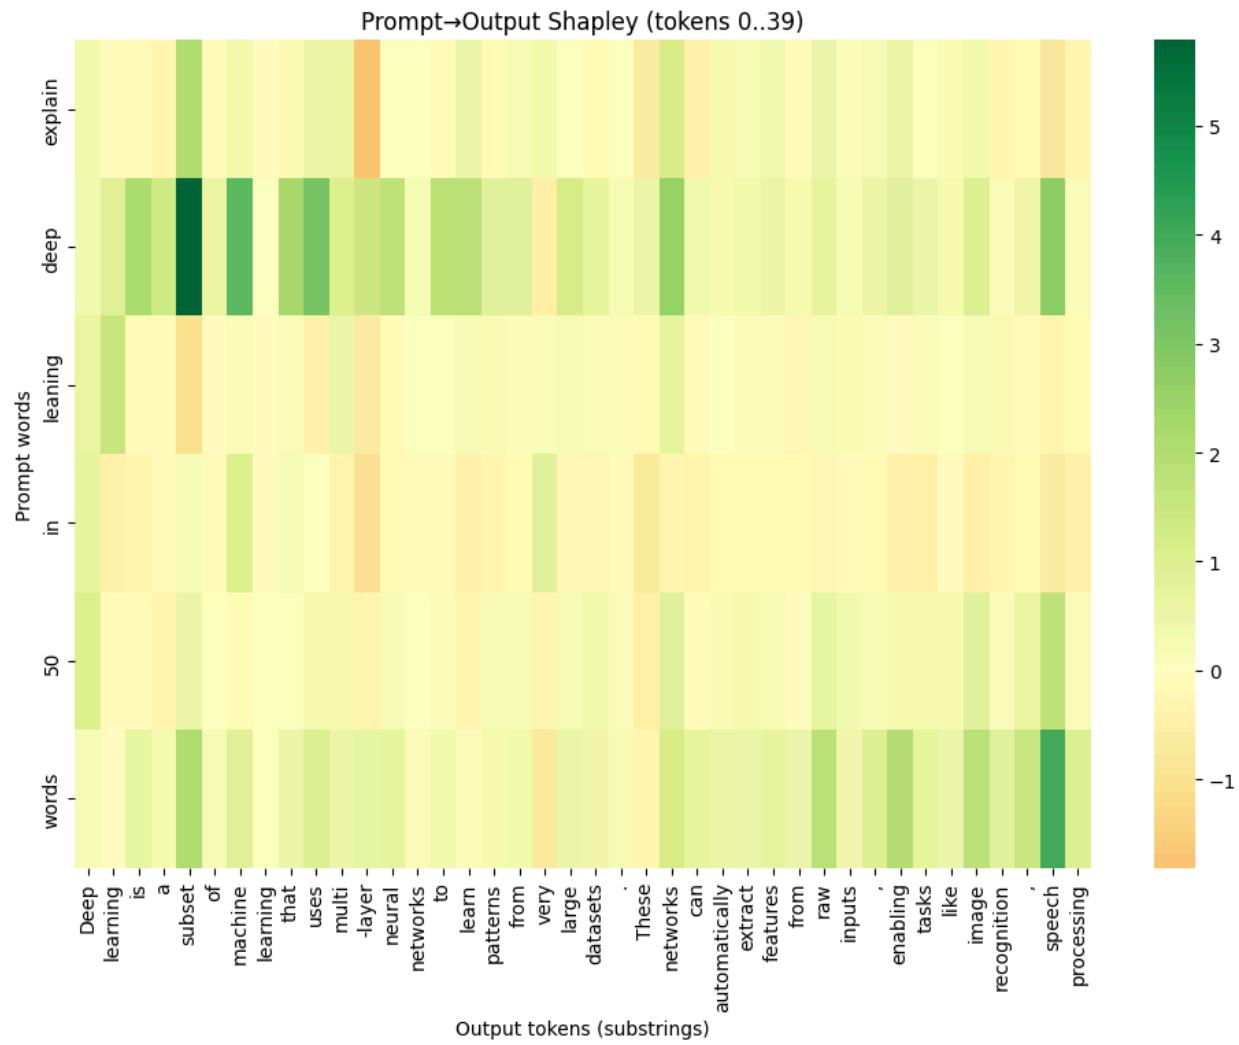
```

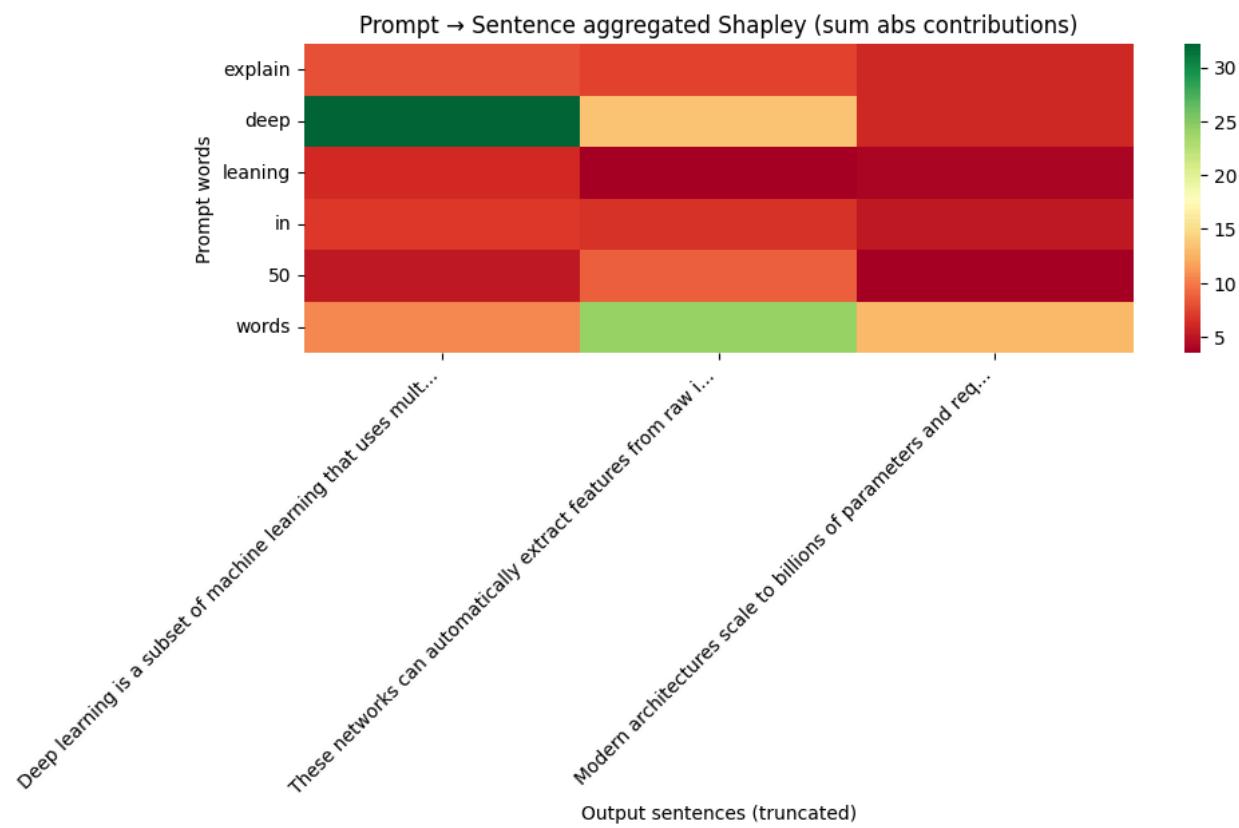
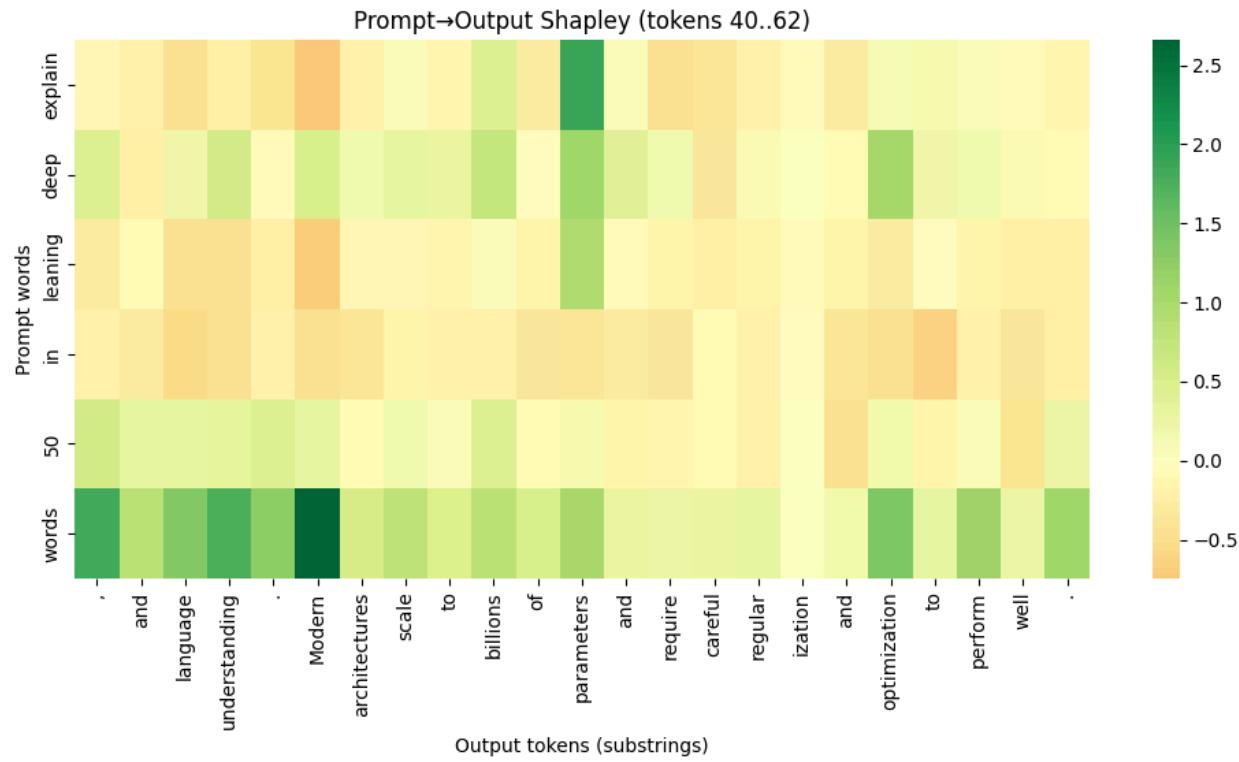
# -----
# Top-k prompt words per output token (for first N tokens) - useful when T large
# -----
def topk_prompt_for_each_output(shap_matrix, prompt_labels, output_labels, t,
                                topk=3, limit_tokens=60):
    n, T = shap_matrix.shape
    for j in range(min(T, limit_tokens)):
        col = shap_matrix[:, j]
        # show absolute magnitude ranking (in logprob units)
        idxs = np.argsort(-np.abs(col))[:topk]
        print(f"Output token [{j}] '{output_labels[j]}' top {topk} prompt influence
              s:")
        for rank, i in enumerate(idxs, 1):
            print(f" {prompt_labels[i]} → contribution = {col[i]: .6f}")
    print("...done (showed first", min(T, limit_tokens), "output tokens.)")

topk_prompt_for_each_output(shap_raw_matrix, prompt_labels, output_token_
strs, topk=3, limit_tokens=40)

```

```
Output token [0] 'Deep' top 3 prompt influences:  
50      -> contribution = 1.091667  
in      -> contribution = 0.709896  
leaning -> contribution = 0.598958  
  
Output token [1] 'learning' top 3 prompt influences:  
leaning -> contribution = 1.487337  
deep    -> contribution = 0.883431  
in     -> contribution = -0.459375  
  
Output token [2] 'is' top 3 prompt influences:  
deep    -> contribution = 2.144141  
words   -> contribution = 0.671257  
in     -> contribution = -0.319352  
  
Output token [3] 'a' top 3 prompt influences:  
deep    -> contribution = 1.377311  
50      -> contribution = -0.347587  
explain -> contribution = -0.291854  
  
Output token [4] 'subset' top 3 prompt influences:  
deep    -> contribution = 5.785514  
explain -> contribution = 2.023413  
words   -> contribution = 2.005501  
...  
Output token [27] 'features' top 3 prompt influences:  
words   -> contribution = 0.746680  
deep    -> contribution = 0.537240  
explain -> contribution = 0.355599
```





Impact of single word on sentences

```
# =====
# SHAPLEY VISUALIZATION (FAST DEMO)
# =====
import torch
import numpy as np
import random
from tqdm import tqdm
import seaborn as sns
import matplotlib.pyplot as plt
from unsloth import FastLanguageModel
from transformers import TextStreamer

# -----
# STEP 1: Load Unsloth Model (lightweight demo)
# -----
# model_name = "unsloth/gpt-oss-20b-unsloth-bnb-4bit"
# model, tokenizer = FastLanguageModel.from_pretrained(
#     model_name=model_name,
#     dtype=None,
#     max_seq_length=4096,
#     load_in_4bit=False,
#     full_finetuning=False,
# )
device = "cuda" if torch.cuda.is_available() else "cpu"
#model = model.to(device)
model.eval()

# -----
# STEP 2: Prompt & hardcoded model output
# -----
prompt = (
```

```

    "explain the role of gen ai in construction industry in 8 sentences"
)
prompt_words = prompt.split()

# Hardcoded output (8-10 sentences)
output_text = (
    "Generative AI (Gen AI) is transforming the construction industry by automating design, planning, and project management tasks. It can generate optimized building designs based on constraints like cost, materials, and sustainability goals. AI-driven simulations help predict structural performance and detect potential issues before construction begins, reducing errors and rework. In project management, Gen AI analyzes schedules, budgets, and workforce data to predict delays and suggest corrective actions. It also assists in creating realistic visualizations and digital twins, improving communication between architects, engineers, and clients. On construction sites, AI-powered tools enhance safety by identifying hazards through real-time image and sensor analysis. Additionally, Gen AI supports material optimization and waste reduction by recommending efficient resource usage. Overall, it drives productivity, sustainability, and innovation across the entire construction lifecycle."
)

# -----
# STEP 3: Compute per-token logprobs
# -----

def compute_logprobs_for_prompt(prompt_text: str, target_text: str):
    inputs = tokenizer.apply_chat_template(
        [{"role": "user", "content": prompt_text}],
        add_generation_prompt=True,
        return_tensors="pt",
        return_dict=True
    ).to(device)

    tgt_ids = tokenizer(target_text).input_ids
    all_input_ids = torch.cat([inputs.input_ids, torch.tensor([tgt_ids], device=device)], dim=1)

```



```

    return arr

for _ in tqdm(range(num_permutations), desc="Shapley perms"):
    perm = list(range(n))
    random.shuffle(perm)
    S = []
    for idx in perm:
        S_sorted = tuple(sorted(S))
        S_with = tuple(sorted(S + [idx]))
        vS = v_of_subset(S_sorted)
        vSi = v_of_subset(S_with)
        marginal = vSi - vS
        accumulator[idx] += marginal
        S.append(idx)

    return accumulator / float(num_permutations)

print("\nComputing Shapley contributions...")
shap_matrix = approximate_shapley_sentence(prompt_words, output_text, compute_logprobs_for_prompt)

# -----
# STEP 5: Aggregate by output sentences
# -----
def aggregate_by_sentences(shap_matrix, output_text, tokenizer, prompt_labels, agg="sum_abs"):
    sentences = [s.strip() for s in output_text.split(".") if s.strip()]
    token_counts = [len(tokenizer(s).input_ids) for s in sentences]
    idx = 0
    ranges = []
    for c in token_counts:
        ranges.append((idx, idx + c))
        idx += c

    n_words, _ = shap_matrix.shape
    n_sents = len(ranges)

```

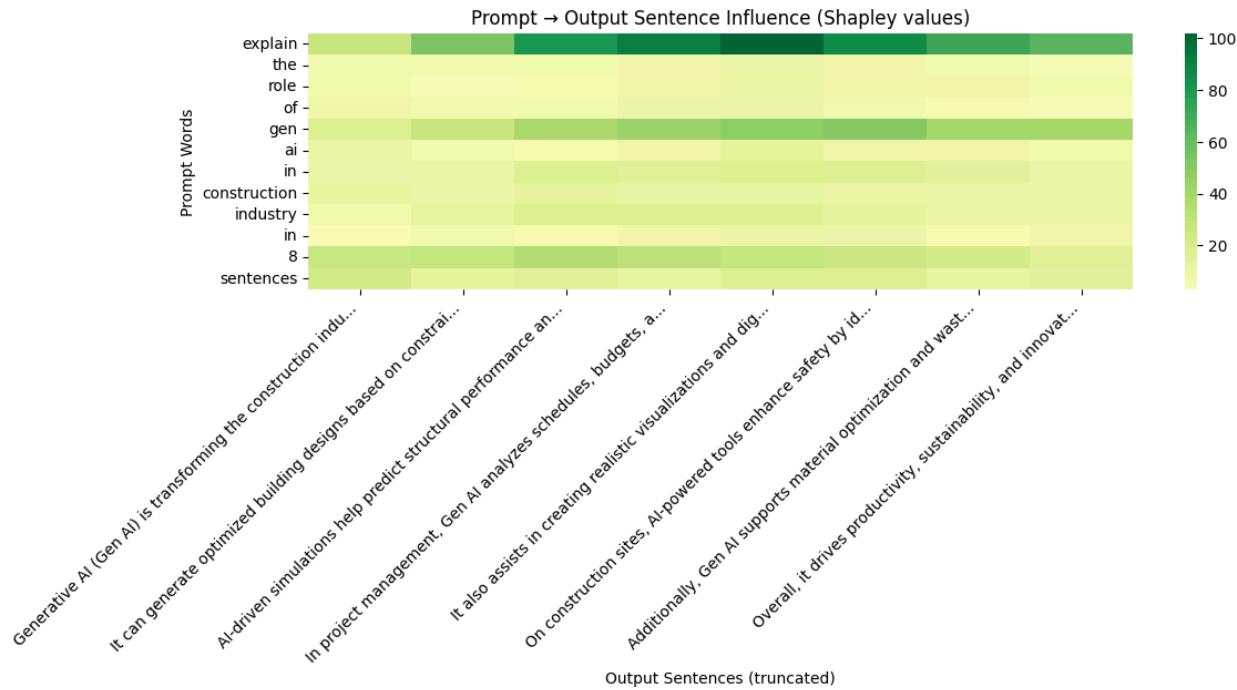
```

agg_mat = np.zeros((n_words, n_sents))
for si, (a, b) in enumerate(ranges):
    if agg == "sum_abs":
        agg_mat[:, si] = np.abs(shap_matrix[:, a:b]).sum(axis=1)
    else:
        agg_mat[:, si] = shap_matrix[:, a:b].sum(axis=1)
return agg_mat, sentences

agg_mat, sents = aggregate_by_sentences(shap_matrix, output_text, tokenize
r, prompt_words)

# -----
# STEP 6: Visualize as heatmap
# -----
plt.figure(figsize=(12, max(3, len(prompt_words) * 0.25)))
sns.heatmap(
    agg_mat,
    xticklabels=[(s[:60] + "...") if len(s) > 60 else s for s in sents],
    yticklabels=prompt_words,
    cmap="RdYIGn",
    center=0.0
)
plt.title("Prompt → Output Sentence Influence (Shapley values)")
plt.xlabel("Output Sentences (truncated)")
plt.ylabel("Prompt Words")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()

```



Recap of What This Code Does (High-Level)

Here's what happens step by step — explained like a human, not a manual:

1. **Prompt setup** — you define the question you want to test.
2. **Hardcoded output** — you use a fixed response (so we don't waste GPU cycles generating).
3. **Log-probability computation** — you run the model on (prompt + output) and get the log probability for each token of the output.
4. **Shapley estimation** — you test what happens when you remove or keep subsets of prompt words, and see how much each word's absence changes output token probabilities.
5. **Aggregate by sentence** — you group tokens into sentences and compute the *total effect* of each prompt word on each output sentence.
6. **Heatmap visualization** — this gives a "Prompt word → Output sentence" influence matrix.

So it answers:

⭐ “If I remove or change this word in the prompt, which sentences in the output are most affected?”